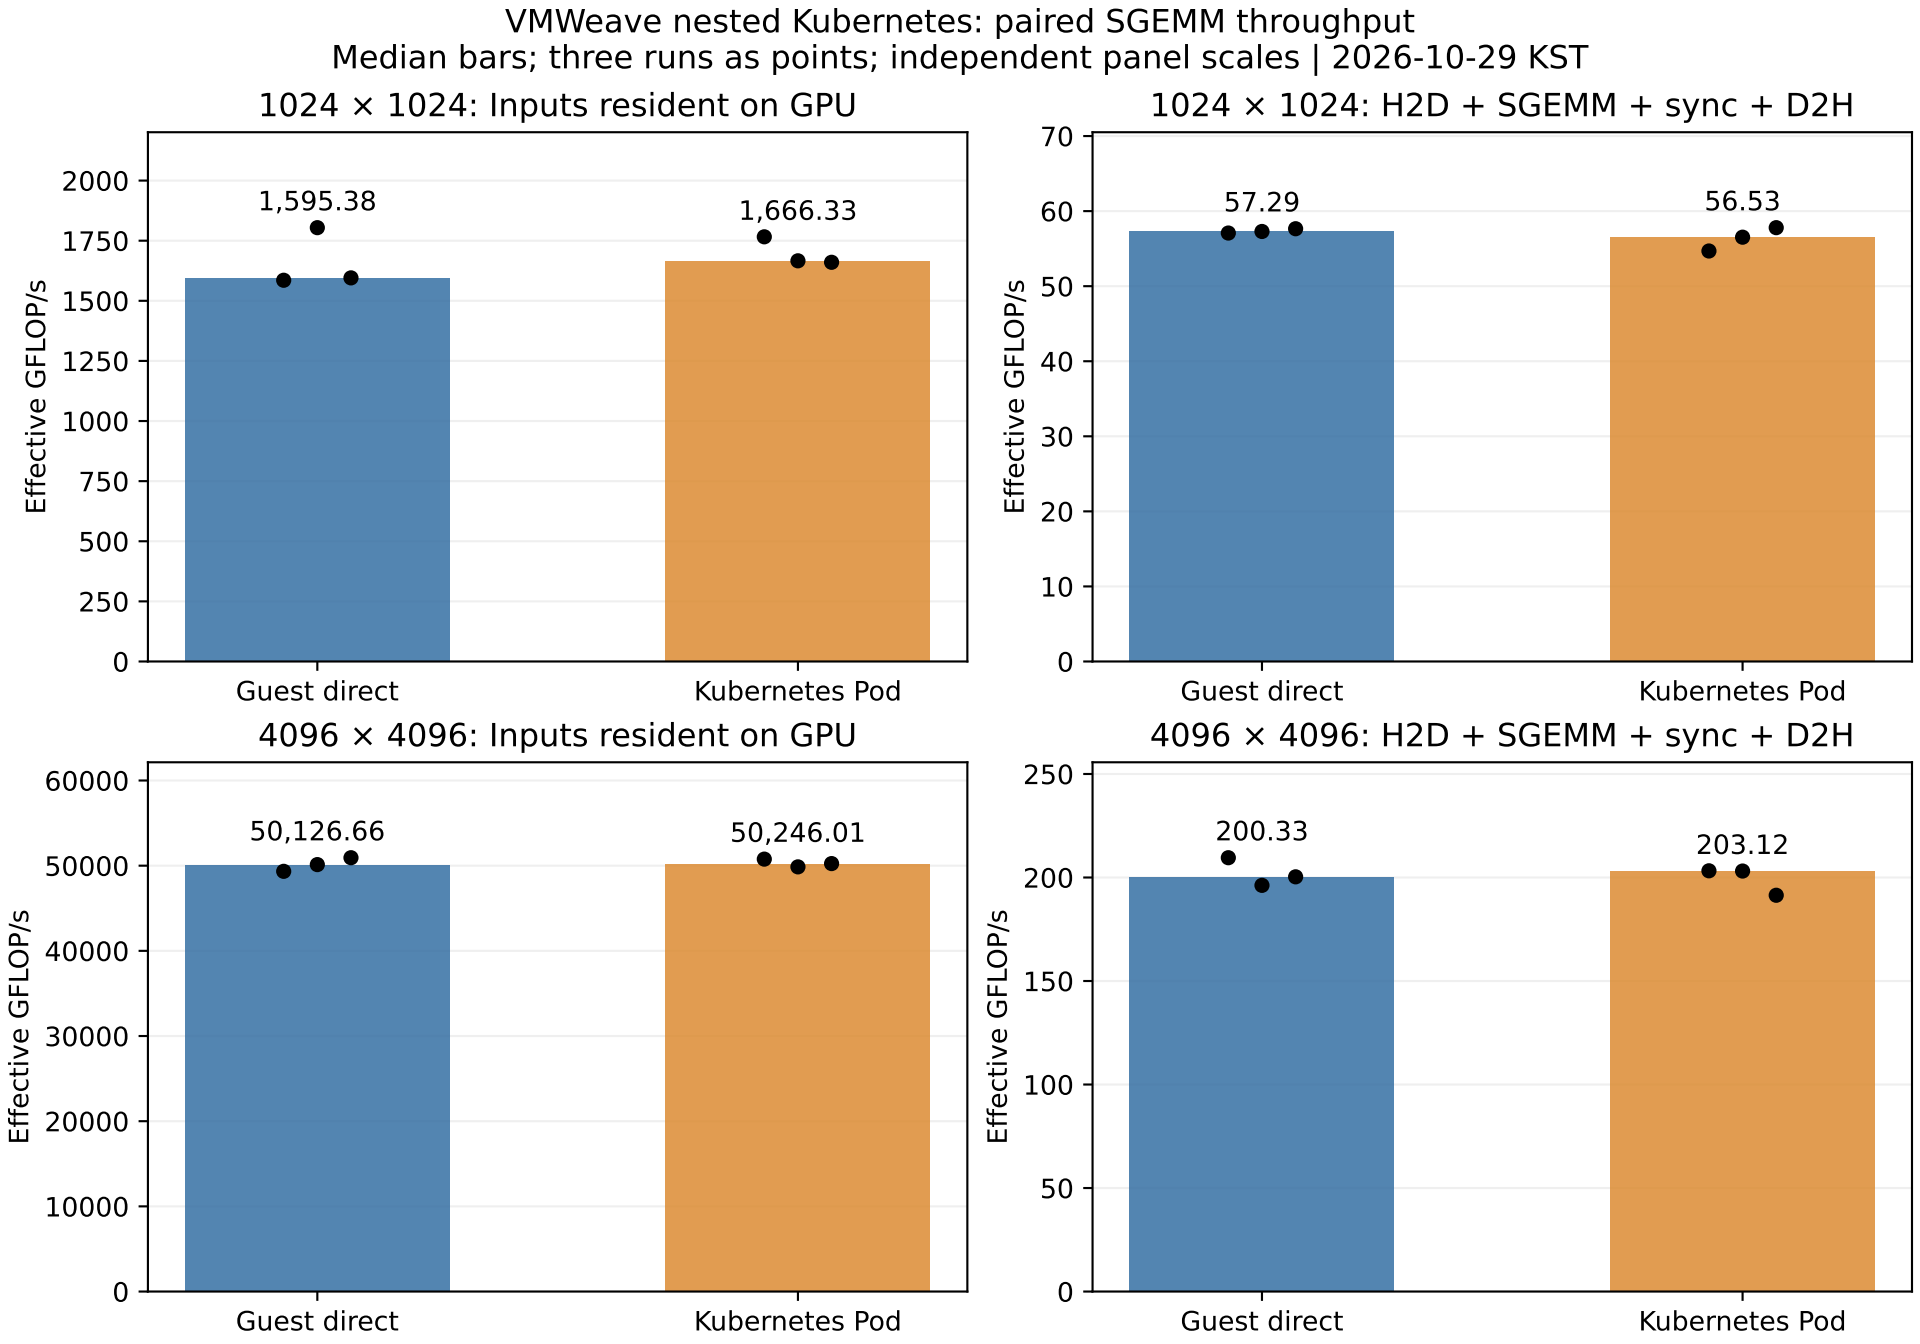
VMWeave nested Kubernetes: paired SGEMM throughput
Median bars; three runs as points; independent panel scales | 2026-10-29 KST
1024 × 1024: Inputs resident on GPU
Effective GFLOP/s
1,595.38
1,666.33
Guest direct
Kubernetes Pod
1024 × 1024: H2D + SGEMM + sync + D2H
57.29
56.53
4096 × 4096: Inputs resident on GPU
50,126.66
50,246.01
4096 × 4096: H2D + SGEMM + sync + D2H
200.33
203.12