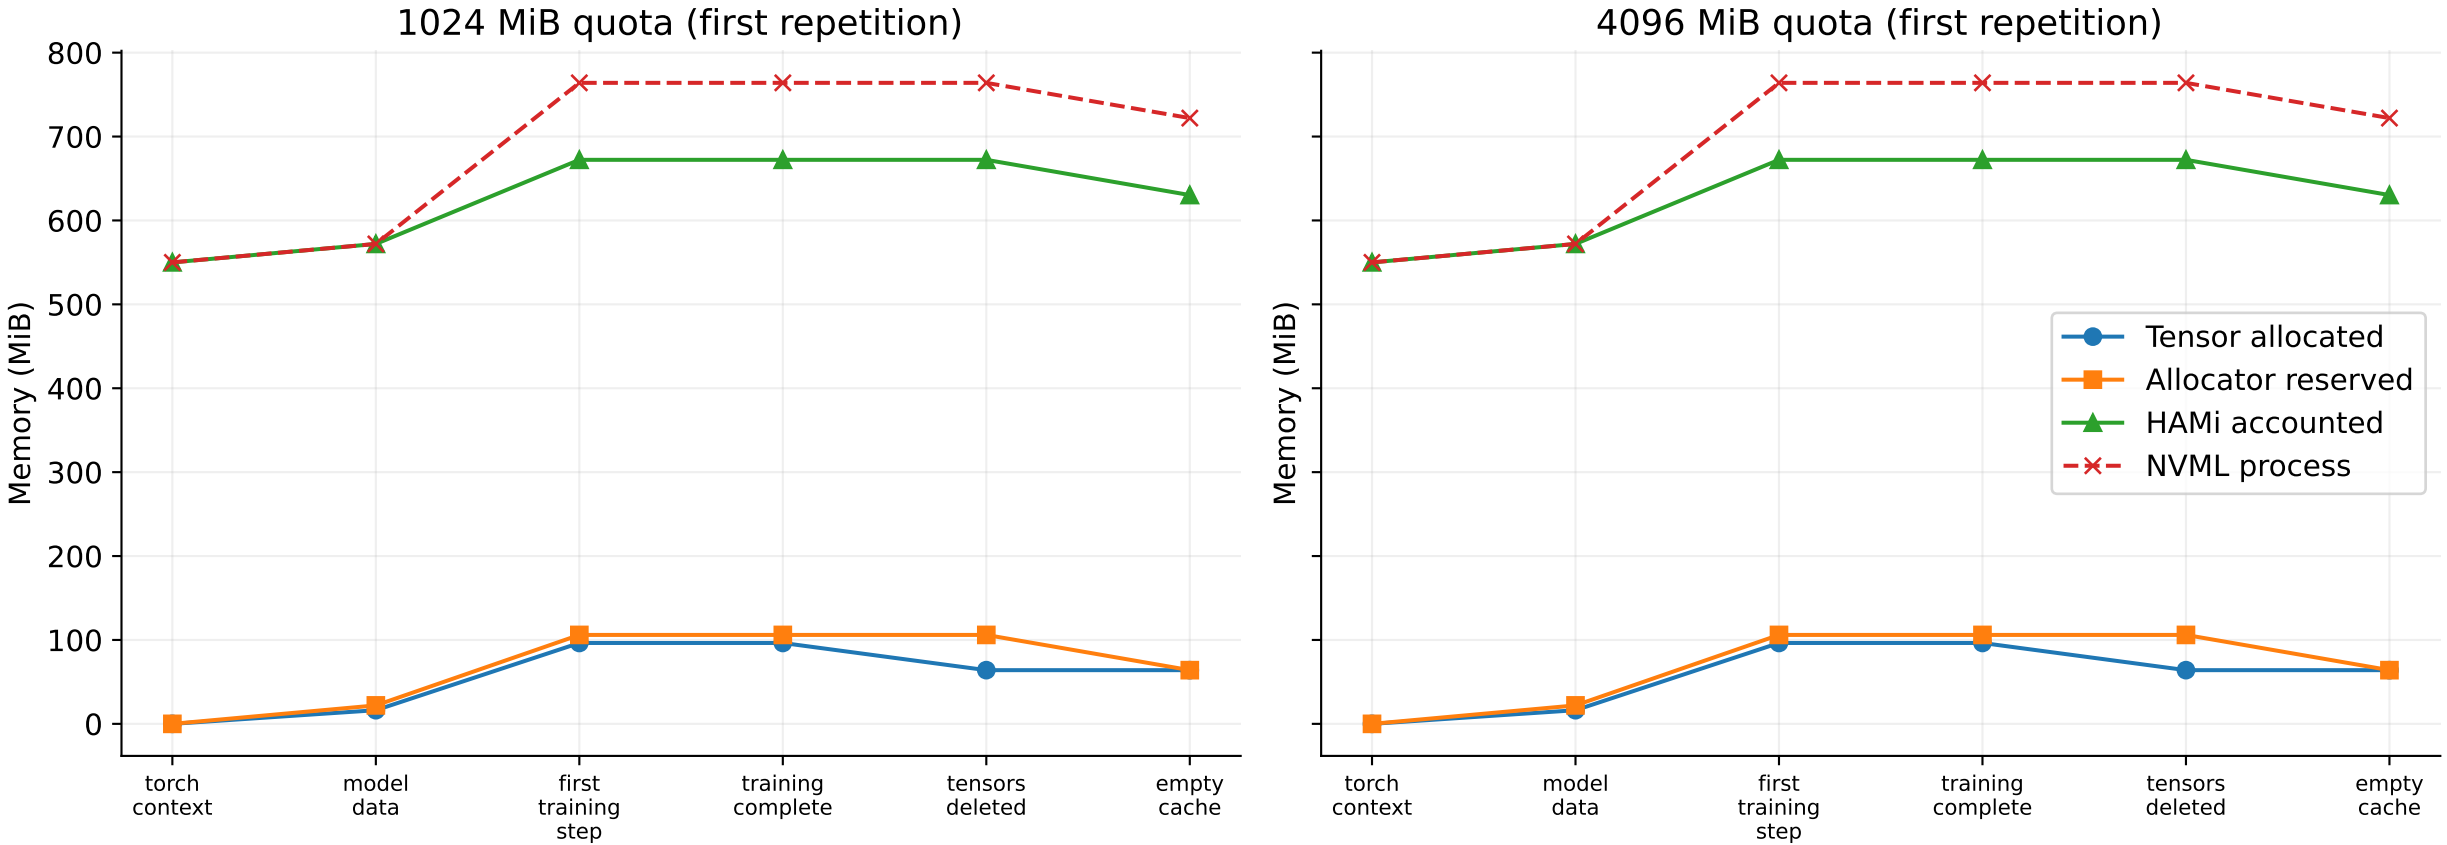
PyTorch allocator metrics do not include all GPU process memory
1024 MiB quota (first repetition)
4096 MiB quota (first repetition)
Memory (MiB)
800
700
600
500
400
300
200
100
0
torch context
model data
first training step
training complete
tensors deleted
empty cache
Tensor allocated
Allocator reserved
HAMi accounted
NVML process